

# Homework 4

## Text Classification with CNN

Welcome to Homework 4!

The homework contains several tasks. You can find the amount of points that you get for the correct solution in the task header. Maximum amount of points for each homework is *four*.

The **grading** for each task is the following:

- correct answer - **full points**
- incorrect solution or solution resulting in the incorrect output - **half points**
- no answer or completely wrong solution - **no points**

Even if you don't know how to solve the task, we encourage you to write down your thoughts and progress and try to address the issues that stop you from completing the task.

When working on the written tasks, try to make your answers short and accurate. Most of the times, it is possible to answer the question in 1-3 sentences.

When writing code, make it readable. Choose appropriate names for your variables (a = 'cat' - not good, word = 'cat' - good). Avoid constructing lines of code longer than 100 characters (79 characters is ideal). If needed, provide the commentaries for your code, however, a good code should be easily readable without them.)

Finally, all your answers should be written only by yourself. If you copy them from other sources it will be considered as an academic fraud. You can discuss the tasks with your classmates but each solution must be individual.

**\*\*Important!\*\* before sending your solution, do the Kernel -> Restart & Run All to ensure that all your code works.**

```
In [31]: !pip install datasets torchmetrics --quiet

# [Output: Progress bars for datasets and torchmetrics installation]
```

```
In [41]: import torch
import torch.nn as nn
from datasets import load_dataset
from torch.utils.data import Dataset, DataLoader
from torchmetrics.functional import f1_score, accuracy
import json
from tqdm.notebook import tqdm
```

In this homework, you are going to work again with [Stanford Semantic Treebank](#).

Only this time, we are going to split the labels into five classes, instead of two. This way, we will do a multi-class classification.

Run the cell below to load the dataset.

```
In [51]: sst = load_dataset("sst", "default")

# [Output: Downloading and preparing dataset sst/default (downloads: 6.83 MiB, generated: 3.73 MiB, post-processed: Un
known size, total: 10.56 MiB) to /root/.cache/huggingface/datasets/sst/default/1.0.0/b8a7889ef01c5d3ae8c37
9b84cc080f8aad3ac2bc538701cbe0ac6416fb76ff...]

Dataset sst downloaded and prepared to /root/.cache/huggingface/datasets/sst/default/1.0.0/b8a7889ef01c5d3
ae8c379b84cc080f8aad3ac2bc538701cbe0ac6416fb76ff. Subsequent calls will reuse this data.

Download the pretrained GloVe 6B word embeddings.

If you are on Windows, just copy the URL into your browser, download the ZIP file and unpack it into the same folder as this
notebook.

In [61]: !wget https://nlp.stanford.edu/data/glove.6B.zip
!unzip glove.6B.zip

--2022-03-08 06:02:59-- https://nlp.stanford.edu/data/glove.6B.zip
Resolving nlp.stanford.edu [nlp.stanford.edu]... 171.64.67.140
Connecting to nlp.stanford.edu [nlp.stanford.edu|171.64.67.140]:443... connected.
HTTP request sent, awaiting response... 301 Moved Permanently
Location: http://downloads.cs.stanford.edu/nlp/data/glove.6B.zip [following]
--2022-03-08 06:02:59-- http://downloads.cs.stanford.edu/nlp/data/glove.6B.zip
Resolving downloads.cs.stanford.edu [downloads.cs.stanford.edu]... 171.64.64.22
Connecting to downloads.cs.stanford.edu [downloads.cs.stanford.edu|171.64.64.22]:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 862182613 (822M) [application/zip]
Saving to: 'glove.6B.zip'

glove.6B.zip           100% [=====] 822.24M  5.06MB/s   in 2m 40s

2022-03-08 06:05:39 [5114 MB/s] - 'glove.6B.zip' saved [862182613/862182613]

Archive: glove.6B.zip
  inflating: glove.6B.50d.txt
  inflating: glove.6B.100d.txt
  inflating: glove.6B.200d.txt
  inflating: glove.6B.300d.txt

In [71]: # Load the embeddings into the memory
glove_path = 'glove.6B.300d.txt'
glove_vecs = []
idx2token = {}

with open(glove_path, encoding='utf-8') as f:
    for line in tqdm(f):
        line = line.strip().split()
        word = line[0]
        vec = [float(x) for x in line[1:]]
        glove_vecs.append(vec)
        idx2token.append(word)

# Convert to tensor
glove_vecs = torch.tensor(glove_vecs)

# Put zero vector for padding and mean for unknown
glove_vecs = torch.vstack([
    torch.zeros(1, glove_vecs.size(1)),
    torch.mean(glove_vecs, dim=0).unsqueeze(0),
    glove_vecs,
])

# Save the embeddings in Pytorch format
torch.save(glove_vecs, 'glove.6B.300d.pt')

# Add special pad and unk tokens to the vocab
PAD = 'pad'
PAD_ID = 0
UNK = 'unk'
UNK_ID = 1

idx2token = [PAD, UNK] + idx2token

# Save the vocab
json.dump(idx2token, open('idx2token.json', 'w', encoding='utf-8'))
```

Uncomment the cell below if you want to load the saved embeddings and vocabulary

```
In [81]: # PAD = 'pad'
# PAD_ID = 0
# UNK = 'unk'
# UNK_ID = 1

# glove_vecs = torch.load('glove.6B.300d.pt')
# idx2token = json.load(open('idx2token.json', encoding='utf-8'))

In [91]: # Create a reverse vocab
token2idx = {token: idx for idx, token in enumerate(idx2token)}

In [101]: # Set the device (GPU or CPU)
device = torch.device('cuda') if torch.cuda.is_available() else torch.device('cpu')
device

Out[10]: device(type='cuda')
```

## Task 1. Build the dataset (1 point)

The SST dataset has a sentiment label for each sentence. This label ranges from 0 to 1, 0 being very negative and 1 being very positive. During the practice session, we split the labels into a binary format, i.e. all the labels lower than 0.5 became 0 and greater than 0.5 became 1.

This time, you will need to split them into five categories that have the following ranges: (0, 0.2], (0.2, 0.4], (0.4, 0.6], (0.6, 0.8], (0.8, 1]. They should be transformed into the labels 0, 1, 2, 3, 4 respectively.

In the end, you should have five labels:

- 0 (very negative)
- 1 (negative)
- 2 (neutral)
- 3 (positive)
- 4 (very positive)

```
In [111]: class SSTDataset(Dataset):
def __init__(self, data, token2idx, max_seq_len=100, device=torch.device('cpu')):
    super().__init__()
    self.max_seq_len = max_seq_len
    self.device = device

    self.texts = []
    self.labels = []
    for item in data:
        label = item['label']
        tokens = token2idx.get(token.lower(), UNK_ID) for token in item['tokens'].split(' ')
        tokens = torch.tensor(tokens, dtype=torch.long, device=self.device)
        self.texts.append(tokens)

        # Transform the continuous label into five classes and add it to the self.labels list
        ## YOUR CODE STARTS HERE
        breakpoints=[0,0.2,0.4,0.6,0.8]
        from bisect import bisect
        self.labels.append(bisect(breakpoints, label)-1)
        ## YOUR CODE ENDS HERE

    def __getitem__(self, idx):
        padded_text = torch.zeros(self.max_seq_len, dtype=torch.long, device=self.device)
        text = self.texts[idx][:self.max_seq_len]
        padded_text[:text.size(0)] = text
        return padded_text, self.labels[idx]

    def __len__(self):
        return len(self.texts)

Load the datasets.

In [121]: train_dataset = SSTDataset(sst['train'], token2idx, max_seq_len=52, device=device)
validation_dataset = SSTDataset(sst['validation'], token2idx, max_seq_len=52, device=device)
test_dataset = SSTDataset(sst['test'], token2idx, max_seq_len=52, device=device)

In [131]: len(train_dataset), len(validation_dataset), len(test_dataset)

Out[13]: (8544, 1101, 2210)

In [141]: train_dataloader = DataLoader(train_dataset, batch_size=50)
validation_dataloader = DataLoader(validation_dataset, batch_size=50)
test_dataloader = DataLoader(test_dataset, batch_size=50)
```

## Task 2. Modify the Model (1 point)

Since now we have five classes instead of two, you need to modify the final layer of the model to have five outputs (ref. [nn.Linear](#)).

Also, we are going to finetune the pretrained embeddings this time, so you need to "unfreeze" them (ref. [nn.Embedding](#)).

This task doesn't have ## YOUR CODE STARTS HERE field, you have to find the parameters and modify them yourself.

```
In [151]: class SSTClassificationModel(nn.Module):
def __init__(self, pretrained_emb, num_filters, kernel_sizes):
    super().__init__()
    self.embedding = nn.Embedding.from_pretrained(pretrained_emb, padding_idx=PAD_ID, freeze=False)
    emb_size = self.embedding.embedding_dim

    self.convs = nn.ModuleList(
        [
            nn.Conv1d(in_channels=emb_size, out_channels=num_filters, kernel_size=kernel_size)
            for kernel_size in kernel_sizes
        ]
    )
    self.linear_out = nn.Linear(in_features=num_filters * len(kernel_sizes), out_features=5)
    self.drop = nn.Dropout(0.5)

    def forward(self, x):
        # x size is [batch x seq_len]
        x = self.embedding(x) # [batch x seq_len x emb_dim]
        x = x.permute(0, 2, 1) # [batch x emb_dim x seq_len]
        xs = [torch.relu(conv(x)) for conv in self.convs] # [batch x num_filters x conv_seq_len] x num_ker
        xs = [torch.nn.functional.max_pool1d(x, x.size(2)).squeeze(2) for x in xs] # [batch x num_filters]
        x = torch.cat(xs, dim=1) # [batch x num_filters * num_kernels]
        x = self.drop(x)
        x = self.linear_out(x) # [batch x 2]
        return x

Model parameters. Feel free to modify them if you'd like to.

In [161]: num_filters = 100
kernel_sizes = [3, 4, 5]
lr = 1e-3

num_iters = 100

Initialize the model.

In [171]: model = SSTClassificationModel(glove_vecs, num_filters, kernel_sizes)
model = model.to(device)

In [181]: print(model)

SSTClassificationModel(
  (emb): Embedding(400002, 300, padding_idx=0)
  (convs): ModuleList(
    (0): Conv1d(300, 100, kernel_size=(3,), stride=(1,))
    (1): Conv1d(300, 100, kernel_size=(4,), stride=(1,))
    (2): Conv1d(300, 100, kernel_size=(5,), stride=(1,))
  )
  (linear_out): Linear(in_features=300, out_features=5, bias=True)
  (drop): Dropout(p=0.5, inplace=False)
)
```

Initialize the loss and optimizer. Feel free to use different optimizer if you'd like to.

```
In [191]: loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=lr)

Train the model.

This may take around 25-30 minutes!
```

In the end, you should have around 0.40 accuracy. If the training takes too long or you have time constraints, feel free to reduce the number of epochs.

```
In [201]: best_accuracy = 0.0
for i in range(num_iters):
    current_loss = 0
    model.train()
    for texts, labels in train_dataloader:
        model.zero_grad()
        preds = model(texts)
        labels = torch.tensor(labels).long().to(device)
        loss = loss_fn(preds, labels)
        current_loss += loss
        loss.backward()
        optimizer.step()

    avg_train_loss = current_loss.item() / len(train_dataloader)
    current_loss = 0
    current_acc = 0
    for texts, labels in validation_dataloader:
        with torch.no_grad():
            preds = model(texts)
            labels = torch.tensor(labels).long().to(device)
            loss = loss_fn(preds, labels)
            preds = torch.argmax(torch.log_softmax(preds, dim=1), dim=1)
            acc = accuracy(preds, labels)
            current_loss += loss
            current_acc += acc
    avg_val_loss = current_loss.item() / len(validation_dataloader)
    avg_val_acc = current_acc.item() / len(validation_dataloader)

    if avg_val_acc > best_accuracy:
        print(f'Accuracy increased [{best_accuracy:.4f} -> {avg_val_acc:.4f}]. Saving the model...')
        best_accuracy = avg_val_acc
        torch.save(model, 'model_best.pt')

    print(f'Epoch: {i}\tTrain loss: {avg_train_loss:.4f}\tVal loss: {avg_val_loss:.4f}\tVal acc: {avg_val_acc:.4f}')

./usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:8: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor.clone().detach() or sourceTensor.clone().detach().requires_grad_(True), rather than torch.tensor(sourceTensor).

./usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:21: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor.clone().detach() or sourceTensor.clone().detach().requires_grad_(True), rather than torch.tensor(sourceTensor).

Accuracy increased [0.0000 -> 0.2974]. Saving the model...
Epoch: 0      Train loss: 1.2682      Val loss: 1.7318      Val acc: 0.2974
Epoch: 1      Train loss: 1.2806      Val loss: 1.7097      Val acc: 0.2496
Epoch: 2      Train loss: 1.1406      Val loss: 1.6654      Val acc: 0.2687
Accuracy increased [0.2974 -> 0.3113]. Saving the model...
Epoch: 3      Train loss: 0.9019      Val loss: 1.7026      Val acc: 0.3113
Accuracy increased [0.3113 -> 0.3365]. Saving the model...
Epoch: 4      Train loss: 0.6330      Val loss: 1.7822      Val acc: 0.3365
Accuracy increased [0.3365 -> 0.3391]. Saving the model...
Epoch: 5      Train loss: 0.0194      Val loss: 2.0280      Val acc: 0.3391
Accuracy increased [0.3391 -> 0.3696]. Saving the model...
Epoch: 6      Train loss: 0.2690      Val loss: 2.4153      Val acc: 0.3696
Accuracy increased [0.3696 -> 0.3739]. Saving the model...
Epoch: 7      Train loss: 0.1948      Val loss: 2.6628      Val acc: 0.3739
Epoch: 8      Train loss: 0.1462      Val loss: 2.7679      Val acc: 0.3739
Epoch: 9      Train loss: 0.1129      Val loss: 3.0582      Val acc: 0.3678
Epoch: 10     Train loss: 0.0886      Val loss: 3.5034      Val acc: 0.3574
Epoch: 11     Train loss: 0.0735      Val loss: 3.7344      Val acc: 0.3591
Epoch: 12     Train loss: 0.0622      Val loss: 4.0025      Val acc: 0.3522
Epoch: 13     Train loss: 0.0516      Val loss: 4.0735      Val acc: 0.3487
Epoch: 14     Train loss: 0.0453      Val loss: 4.3162      Val acc: 0.3626
Epoch: 15     Train loss: 0.0441      Val loss: 4.4438      Val acc: 0.3626
Epoch: 16     Train loss: 0.0376      Val loss: 4.6097      Val acc: 0.3722
Epoch: 17     Train loss: 0.0359      Val loss: 4.7991      Val acc: 0.3722
Epoch: 18     Train loss: 0.0194      Val loss: 4.8185      Val acc: 0.3357
Epoch: 19     Train loss: 0.0310      Val loss: 5.0603      Val acc: 0.3617
Epoch: 20     Train loss: 0.0288      Val loss: 5.1372      Val acc: 0.3783
Epoch: 21     Train loss: 0.0264      Val loss: 5.3437      Val acc: 0.3200
Epoch: 22     Train loss: 0.0259      Val loss: 5.2117      Val acc: 0.3287
Epoch: 23     Train loss: 0.0116      Val loss: 5.1076      Val acc: 0.3339
Epoch: 24     Train loss: 0.0270      Val loss: 5.8346      Val acc: 0.3348
Epoch: 25     Train loss: 0.0244      Val loss: 5.7714      Val acc: 0.3304
Epoch: 26     Train loss: 0.0310      Val loss: 5.6305      Val acc: 0.3391
Accuracy increased [0.3739 -> 0.3783]. Saving the model...
Epoch: 27     Train loss: 0.0254      Val loss: 5.9364      Val acc: 0.3783
Epoch: 28     Train loss: 0.0232      Val loss: 6.1860      Val acc: 0.3304
Epoch: 29     Train loss: 0.0208      Val loss: 6.2371      Val acc: 0.3322
Epoch: 30     Train loss: 0.0234      Val loss: 6.6982      Val acc: 0.3339
Epoch: 31     Train loss: 0.0244      Val loss: 6.8243      Val acc: 0.3357
Epoch: 32     Train loss: 0.0218      Val loss: 6.6238      Val acc: 0.3435
Epoch: 33     Train loss: 0.0278      Val loss: 6.8556      Val acc: 0.3383
Epoch: 34     Train loss: 0.0221      Val loss: 7.0506      Val acc: 0.3304
Epoch: 35     Train loss: 0.0211      Val loss: 7.5199      Val acc: 0.3322
Epoch: 36     Train loss: 0.0170      Val loss: 7.3695      Val acc: 0.3339
Epoch: 37     Train loss: 0.0182      Val loss: 7.5942      Val acc: 0.3426
Epoch: 38     Train loss: 0.0147      Val loss: 7.8464      Val acc: 0.3374
Epoch: 39     Train loss: 0.0213      Val loss: 8.7774      Val acc: 0.3435
Epoch: 40     Train loss: 0.0165      Val loss: 8.1847      Val acc: 0.3421
Epoch: 41     Train loss: 0.0109      Val loss: 8.0018      Val acc: 0.3357
Epoch: 42     Train loss: 0.0180      Val loss: 8.0301      Val acc: 0.3435
Epoch: 43     Train loss: 0.0119      Val loss: 8.0297      Val acc: 0.3452
Epoch: 44     Train loss: 0.0112      Val loss: 8.2545      Val acc: 0.3530
Epoch: 45     Train loss: 0.0156      Val loss: 8.3333      Val acc: 0.3478
Epoch: 46     Train loss: 0.0136      Val loss: 8.4904      Val acc: 0.3426
Accuracy increased [0.3783 -> 0.3835]. Saving the model...
Epoch: 47     Train loss: 0.0128      Val loss: 8.9308      Val acc: 0.3835
Epoch: 48     Train loss: 0.0135      Val loss: 8.5183      Val acc: 0.3357
Epoch: 49     Train loss: 0.0126      Val loss: 8.8779      Val acc: 0.3400
Epoch: 50     Train loss: 0.0196      Val loss: 12.6882      Val acc: 0.3287
Epoch: 51     Train loss: 0.0142      Val loss: 9.4477      Val acc: 0.3339
Epoch: 52     Train loss: 0.0120      Val loss: 9.2859      Val acc: 0.3374
Epoch: 53     Train loss: 0.0219      Val loss: 9.0855      Val acc: 0.3426
Epoch: 54     Train loss: 0.0122      Val loss: 9.2997      Val acc: 0.3426
Epoch: 55     Train loss: 0.0124      Val loss: 9.3660      Val acc: 0.3452
Epoch: 56     Train loss: 0.0119      Val loss: 9.1815      Val acc: 0.3426
Epoch: 57     Train loss: 0.0119      Val loss: 9.2447      Val acc: 0.3400
Epoch: 58     Train loss: 0.0101      Val loss: 9.4155      Val acc: 0.3339
Epoch: 59     Train loss: 0.0120      Val loss: 9.4888      Val acc: 0.3322
Epoch: 60     Train loss: 0.0113      Val loss: 10.0384      Val acc: 0.3387
Epoch: 61     Train loss: 0.0113      Val loss: 10.1186      Val acc: 0.3409
Epoch: 62     Train loss: 0.0100      Val loss: 10.0781      Val acc: 0.3374
Epoch: 63     Train loss: 0.0089      Val loss: 10.0744      Val acc: 0.3435
Epoch: 64     Train loss: 0.0095      Val loss: 10.1380      Val acc: 0.3435
Epoch: 65     Train loss: 0.0099      Val loss: 10.0253      Val acc: 0.3357
Epoch: 66     Train loss: 0.0092      Val loss: 10.0662      Val acc: 0.3383
Accuracy increased [0.3835 -> 0.3922]. Saving the model...
Epoch: 67     Train loss: 0.0104      Val loss: 10.5337      Val acc: 0.3922
Epoch: 68     Train loss: 0.0084      Val loss: 10.3771      Val acc: 0.3400
Epoch: 69     Train loss: 0.0084      Val loss: 10.6182      Val acc: 0.3426
Epoch: 70     Train loss: 0.0082      Val loss: 10.5027      Val acc: 0.3339
Epoch: 71     Train loss: 0.0015      Val loss: 10.6250      Val acc: 0.3374
Epoch: 72     Train loss: 0.0138      Val loss: 10.9179      Val acc: 0.3374
Epoch: 73     Train loss: 0.0091      Val loss: 10.5023      Val acc: 0.3357
Epoch: 74     Train loss: 0.0111      Val loss: 11.1076      Val acc: 0.3417
Epoch: 75     Train loss: 0.0101      Val loss: 11.0918      Val acc: 0.3339
Epoch: 76     Train loss: 0.0086      Val loss: 10.5928      Val acc: 0.3426
Epoch: 77     Train loss: 0.0101      Val loss: 11.1122      Val acc: 0.3400
Epoch: 78     Train loss: 0.0109      Val loss: 11.2948      Val acc: 0.3357
Epoch: 79     Train loss: 0.0105      Val loss: 11.6101      Val acc: 0.3409
Epoch: 80     Train loss: 0.0093      Val loss: 11.6067      Val acc: 0.3452
Epoch: 81     Train loss: 0.0077      Val loss: 11.6515      Val acc: 0.3487
Epoch: 82     Train loss: 0.0084      Val loss: 11.9468      Val acc: 0.3400
Epoch: 83     Train loss: 0.0076      Val loss: 11.9358      Val acc: 0.3339
Epoch: 84     Train loss: 0.0070      Val loss: 12.1196      Val acc: 0.3417
Epoch: 85     Train loss: 0.0068      Val loss: 12.2693      Val acc: 0.3426
Epoch: 86     Train loss: 0.0058      Val loss: 12.4514      Val acc: 0.3348
Epoch: 87     Train loss: 0.0085      Val loss: 12.7323      Val acc: 0.3426
Epoch: 88     Train loss: 0.0076      Val loss: 12.6882      Val acc: 0.3400
Epoch: 89     Train loss: 0.0058      Val loss: 12.7372      Val acc: 0.3409
Epoch: 90     Train loss: 0.0075      Val loss: 13.0063      Val acc: 0.3461
Epoch: 91     Train loss: 0.0110      Val loss: 13.2705      Val acc: 0.3478
Epoch: 92     Train loss: 0.0091      Val loss: 12.8850      Val acc: 0.3496
Epoch: 93     Train loss: 0.0094      Val loss: 12.8107      Val acc: 0.3452
Epoch: 94     Train loss: 0.0092      Val loss: 13.0835      Val acc: 0.3504
Epoch: 95     Train loss: 0.0085      Val loss: 13.2118      Val acc: 0.3523
Epoch: 96     Train loss: 0.0063      Val loss: 13.1689      Val acc: 0.3443
Epoch: 97     Train loss: 0.0074      Val loss: 13.4532      Val acc: 0.3504
Epoch: 98     Train loss: 0.0089      Val loss: 13.4954      Val acc: 0.3467
Epoch: 99     Train loss: 0.0082      Val loss: 13.9221      Val acc: 0.3522

Load the best model.

In [211]: model = torch.load('model_best.pt')

In [221]: model.to(device)

Out[221]: SSTClassificationModel(
  (emb): Embedding(400002, 300, padding_idx=0)
  (convs): ModuleList(
    (0): Conv1d(300, 100, kernel_size=(3,), stride=(1,))
    (1): Conv1d(300, 100, kernel_size=(4,), stride=(1,))
    (2): Conv1d(300, 100, kernel_size=(5,), stride=(1,))
  )
  (linear_out): Linear(in_features=300, out_features=5, bias=True)
  (drop): Dropout(p=0.5, inplace=False)
)
```

Test the model on the test set and save all the labels and predictions.

```
In [231]: current_acc = 0
current_loss = 0
y_true = []
y_pred = []
model.eval()
with torch.no_grad():
    for text, label in test_dataloader:
        label = torch.tensor(label).long().to(device)
        preds = model(text)
        loss = loss_fn(preds, label)
        preds = torch.argmax(torch.log_softmax(preds, dim=1), dim=1)
        acc = accuracy(preds, label)
        current_loss += loss
        current_acc += acc
        y_true.append(label)
        y_pred.append(preds)
avg_test_loss = current_loss.item() / len(test_dataloader)
avg_test_acc = current_acc.item() / len(test_dataloader)

print(f'Test loss: {avg_test_loss:.6f}\tTest acc: {avg_test_acc:.2%}')

y_true = torch.hstack(y_true).cpu().numpy()
y_pred = torch.hstack(y_pred).cpu().numpy()

./usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:8: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor.clone().detach() or sourceTensor.clone().detach().requires_grad_(True), rather than torch.tensor(sourceTensor).

Test loss: 10.209472      Test acc: 38.67%
```

## Task 3. Evaluate the model (1 point)

One way to better understand a multi-class model is to build a **confusion matrix**.

Run the cell below and describe what you see. For example, which classes are the most difficult for the model to differentiate? What one are the easiest? Why do you think it happens?

```
In [241]: from sklearn.metrics import ConfusionMatrixDisplay

class_labels = ['very negative', 'negative', 'neutral', 'positive', 'very positive']
ConfusionMatrixDisplay.from_predictions(y_true, y_pred, display_labels=class_labels, xticks_rotation="vertical")
```

**YOUR ANSWER HERE**

(A): In sentimental meaning term, it can be easily to replace and correct in grammatical fluency(e.g., nice, better ). But when I try unsentimental term (e.g. king, dog) both still have nearly same word but different score.